

Paramount+ — Day-Of Cram Sheet

60-sec intro

- 15 yrs Python/AWS · Disney media/financial backend · \$10M+ daily · batch 8hr→45min · 99.99% uptime
- CGJSoftware Lambda/ECS · 2.5M daily txns · <120ms
- Why Paramount: scalable media platform — ingest, catalog, HA APIs

Disney STAR (payment team)

- Monolith SQL → pitched API + UI → hub-and-spoke (Jira, Airtable, costing)
- Result: month-end payroll **8hr → 45min** in 3 months
- Hulu acquisition: streaming-adjacent context; owned batch/reporting on AWS

AWS quick hits

- **Lambda**: event-driven, spiky · **Fargate**: long-running API · hybrid both
- **Ingest**: S3 → SQS → Lambda → DynamoDB → Step Functions → MediaConvert → CDN
- **Lambda code**: thin handler → service → repo · idempotent on messageId
- **Step Functions**: Retry transient · Catch permanent · SNS alert
- **DynamoDB**: PK CONTENT#id · SK METADATA/EPISODE#n · GSI slug + genre/date

AI tools

- Cursor/Copilot for tests/refactors · 65%→92% coverage · never ship unreviewed

Ask them

- What does media platform team own? · Lambda vs ECS split? · DynamoDB in catalog?

Day-of tip: Open this HTML in Chrome/Safari — faster than PDF. Don't re-render PDF during the call.